

```
In [1]: from google.colab import drive
drive.mount("/content/drive/")
```

Drive already mounted at /content/drive/; to attempt to forcibly remount, call drive.mount("/content/drive/", force_remount=True).

```
In [2]: !wget --header="Host: storage.googleapis.com" --header="User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/90.0.4430.212 Safari/537.36" --header="Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.9" --header="Accept-Language: en-US,en;q=0.9" --header="Referer: https://www.kaggle.com/" "https://storage.googleapis.com/kaggle-data-sets/623289/1111676/bundle/archive.zip?X-Goog-Algorithm=G00G4-RSA-SHA256&X-Goog-Credential=gcp-kaggle-com%40kaggle-161607.iam.gserviceaccount.com%2F20210526%2Fauto%2Fstorage%2Fgoog4_request&X-Goog-Date=20210526T170031Z&X-Goog-Expires=259199&X-Goog-SignedHeaders=host&X-Goog-Signature=356683d4035c5ce3fb1aa29ad63a9696248263723c3c5365249aa55c7f88a91841c92b9a571d67d97f437227d4b7747dc257dab9f42f6f1a8ac7c30933d839bf4d7dab62999a62549f529910e4e1963ff3eb24007268349e07c584d6a44f8696c55a2f0f403ea01cbaa6c912b73648feee1880d4a5c13ee41d376070fd569c0dda2b99bd5115ef08e71e9854a0d8334f6d14769c267af22bc83e05fd5549e176e871d02ef6dbe9f5c9a467c2fa305adbf17bc12a1697b33a7ae43dd50aa2173547c783c666be2a9d8db62feee63384413372590c86c99ba9ef5027394a18c559ca776b48a8fd9b2177ae4d47fd2db67450b69d2bc810857cbe71990820ff2f8" -c -O 'archive.zip'
```

```
--2021-05-28 17:41:48-- https://storage.googleapis.com/kaggle-data-sets/623289/1111676/bundle/archive.zip?X-Goog-Algorithm=G00G4-RSA-SHA256&X-Goog-Credential=gcp-kaggle-com%40kaggle-161607.iam.gserviceaccount.com%2F20210526%2Fauto%2Fstorage%2Fgoog4_request&X-Goog-Date=20210526T170031Z&X-Goog-Expires=259199&X-Goog-SignedHeaders=host&X-Goog-Signature=356683d4035c5ce3fb1aa29ad63a9696248263723c3c5365249aa55c7f88a91841c92b9a571d67d97f437227d4b7747dc257dab9f42f6f1a8ac7c30933d839bf4d7dab62999a62549f529910e4e1963ff3eb24007268349e07c584d6a44f8696c55a2f0f403ea01cbaa6c912b73648feee1880d4a5c13ee41d376070fd569c0dda2b99bd5115ef08e71e9854a0d8334f6d14769c267af22bc83e05fd5549e176e871d02ef6dbe9f5c9a467c2fa305adbf17bc12a1697b33a7ae43dd50aa2173547c783c666be2a9d8db62feee63384413372590c86c99ba9ef5027394a18c559ca776b48a8fd9b2177ae4d47fd2db67450b69d2bc810857cbe71990820ff2f8
```

Resolving storage.googleapis.com (storage.googleapis.com)... 172.217.13.80, 172.217.13.240, 142.251.33.208, ...

Connecting to storage.googleapis.com (storage.googleapis.com)|172.217.13.80|:443... connected.

HTTP request sent, awaiting response... 416 Requested range not satisfiable

The file is already fully retrieved; nothing to do.

```
In [3]: #Extracting the
from zipfile import ZipFile
with ZipFile("archive.zip", "r") as z:
    z.extractall()
    print("done")
```

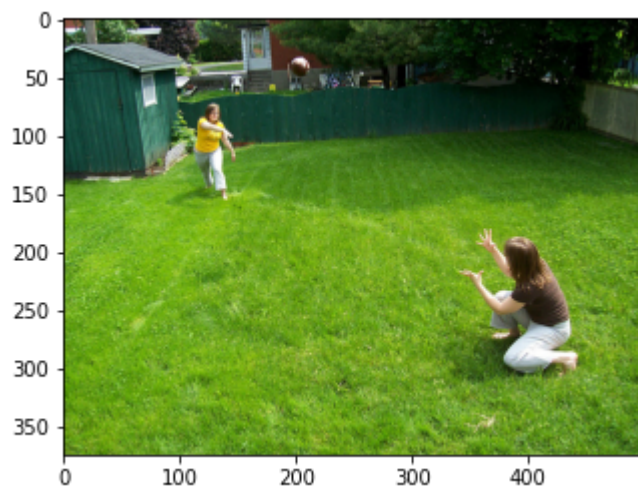
done

Processing The Image

```
In [4]: import matplotlib.pyplot as plt
import cv2
import os
k=os.listdir("Images")
image_list=k[:2000]
image_name=k[:2000]
for i in range(len(image_list)):
    image_list[i]="Images/"+image_list[i]

img=cv2.imread("Images/534200447_b0f3ff02be.jpg")
img=cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
plt.imshow(img)
```

Out[4]: <matplotlib.image.AxesImage at 0x7f8f175853d0>



```
In [ ]: test_images=k[2000:2100]
print(test_images)
for i in range(len(test_images)):
    test_images[i]="Images/"+test_images[i]
```

```
In [43]: import tensorflow as tf
image_model = tf.keras.applications.VGG16(include_top=False, weights='imagenet')
new_input = image_model.input
hidden_layer = image_model.layers[-1].output
image_model = tf.keras.Model(new_input, hidden_layer)
image_model.summary()
```

Model: "model_1"

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	[(None, None, None, 3)]	0
<hr/>		
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
<hr/>		
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
<hr/>		
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
<hr/>		
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
<hr/>		
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
<hr/>		
block2_pool (MaxPooling2D)	(None, None, None, 128)	0
<hr/>		
block3_conv1 (Conv2D)	(None, None, None, 256)	295168
<hr/>		
block3_conv2 (Conv2D)	(None, None, None, 256)	590080
<hr/>		
block3_conv3 (Conv2D)	(None, None, None, 256)	590080
<hr/>		
block3_pool (MaxPooling2D)	(None, None, None, 256)	0
<hr/>		
block4_conv1 (Conv2D)	(None, None, None, 512)	1180160
<hr/>		
block4_conv2 (Conv2D)	(None, None, None, 512)	2359808
<hr/>		
block4_conv3 (Conv2D)	(None, None, None, 512)	2359808
<hr/>		
block4_pool (MaxPooling2D)	(None, None, None, 512)	0
<hr/>		
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
<hr/>		
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
<hr/>		
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
<hr/>		
block5_pool (MaxPooling2D)	(None, None, None, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		
<hr/>		

```
In [ ]: image_model.save("/content/drive/My Drive/CASE Studies/Image captioning/Attention/image_model.h5")
```

```
In [7]: #for that I m writing a function which takes image and coverts it to vector
```

```
def img_2_vec(x):  
    img=cv2.imread(x)  
    img=cv2.resize(img, (224,224))  
    img=img.reshape(1,224,224,3)  
    img=image_model.predict(img)  
    img=img.reshape(49,512)  
    return img  
  
img_vector=[]  
for i in image_list:  
    img_vector.append(img_2_vec(i))
```

```
In [8]: len(img_vector)
```


```
Out[8]: 2000
```

```
In [9]: #saving image_name and its vector in dictionary  
dict_vector={}  
for i in range(len(image_list)):  
    dict_vector[image_list[i].split("/")[1]]=img_vector[i]  
  
len(dict_vector)
```

```
Out[9]: 2000
```

```
In [10]: # image and its feature vector
img=cv2.imread("Images/534200447_b0f3ff02be.jpg")
img=cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
plt.imshow(img)
print("vector=",dict_vector[image_list[5].split("/")[1]])
```

```
vector= [[0.          0.          0.          ... 0.          0.          0.          ]
 [0.          0.          0.          ... 0.          0.          0.          ]
 [0.          0.          0.          ... 0.          3.6347606 0.          ]
 ...
 [0.          0.          0.          ... 0.          0.          0.          ]
 [0.          0.          0.          ... 7.5758348 0.          0.          ]
 [0.          0.          0.          ... 0.          0.          0.          ]]
```



Processing Captions

```
In [11]: # reading the Captions file
with open("captions.txt",'r') as file:
    text=file.readlines()
```

```
In [12]: # creating dictionary which contains image name and its captions
dict_captions={}
for i in text:
    name=i.split(",")[0]
    caption=i.split(",")[1]
    if name not in dict_captions:
        dict_captions[name]=[caption]
    else:
        dict_captions[name].append(caption)
dict_captions["534200447_b0f3ff02be.jpg"]

cap={}
for i in text:
    name=i.split(",")[0]
    caption=i.split(",")[1]
    if name not in cap:
        cap[name]=[caption]
    else:
        cap[name].append(caption)
```

```
In [13]: dict_captions=dict([(key ,dict_captions[key]) for key in image_name])
print(len(dict_captions))
#cap=dict([(key ,cap[key]) for key in image_name])

2000
```

```
In [14]: #Preprocessing the Captions
import re

def preprocess(d):
    for k,v in d.items():
        for i in range(len(v)):
            v[i]=re.sub('[^A-Za-z]+',' ',v[i])
            v[i]=v[i].lower()
            v[i]=v[i].strip()

            v[i]="start "+v[i]+" end"
    return d
dict_captions=preprocess(dict_captions)
cap=preprocess(cap)
dict_captions['2975018306_0e8da316f5.jpg']
```

```
Out[14]: ['start a brown dog is licking its nose end',
'start a brown dog is on the green grass end',
'start a brown dog licking his nose in the middle of some grass end',
'start a brown dog licks his nose while standing in green grass by the woods end',
'start a brown dog with his tongue out looks at the camera end']
```

```
In [15]: #Plotting the image and its captions
import cv2
import matplotlib.pyplot as plt
img=cv2.imread('Images/2975018306_0e8da316f5.jpg')
img=cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
plt.imshow(img)
dict_captions['2975018306_0e8da316f5.jpg']
```

```
Out[15]: ['start a brown dog is licking its nose end',
'start a brown dog is on the green grass end',
'start a brown dog licking his nose in the middle of some grass end',
'start a brown dog licks his nose while standing in green grass by the woods end',
'start a brown dog with his tongue out looks at the camera end']
```



```

In [16]: #Creating Vocabulary for the words and giving them the index/integer value
vocab=set()
for k,v in dict_captions.items():
    for i in range(len(v)):
        for j in v[i].split():
            vocab.add(j)
vocab=sorted(vocab)
vocabulary={}
for i,j in enumerate(vocab,1):
    vocabulary[j]=i

print("The Unique Number of words are ",len(vocab))

#trasforming these index to captions
for k,v in dict_captions.items():
    for i in v:
        t=[]
        for j in i.split():
            t.append(vocabulary[j])

        dict_captions[k][v.index(i)]=t

for i in dict_captions['2975018306_0e8da316f5.jpg']:
    print(i)

```

```

The Unique Number of words are 4329
[3595, 1, 461, 1061, 1887, 2090, 1895, 2404, 1208]
[3595, 1, 461, 1061, 1887, 2458, 3842, 1596, 1586, 1208]
[3595, 1, 461, 1061, 2090, 1764, 2404, 1847, 3842, 2253, 2440, 3471, 1586, 1208]
[3595, 1, 461, 1061, 2091, 1764, 2404, 4229, 3585, 1847, 1596, 1586, 523, 3842, 4283, 1208]
[3595, 1, 461, 1061, 4270, 1764, 3914, 2498, 2146, 145, 3842, 538, 1208]

```

```

In [17]: len(vocab)

```

```

Out[17]: 4329

```

```

In [18]: #finding the maxlenght of captions from train captions
max_len=[]
for k,v in dict_captions.items():
    for i in v:
        max_len.append(len(i))
max_length=max(max_len)
print("The maximum lenght of caption is",max(max_len))

```

```

The maximum lenght of caption is 35

```



```
In [19]: #Creating the dataset
from keras.preprocessing.sequence import pad_sequences
import numpy as np
X = []
y_in = []
for k, v in dict_captions.items():
    for j in v:
        X.append(dict_vector[k])
        in_seq = pad_sequences([j], maxlen=max_length, padding='post', truncating='post')[0]
        y_in.append(in_seq)
X=np.array(X)
y_in=np.array(y_in)
print(X.shape,y_in.shape)

(10000, 49, 512) (10000, 35)
```

```
In [20]: #splitting the according to images names
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(X,y_in,test_size=0.2)
```

```
In [21]: def create_dataset(img_name_train,caption_train):
    dataset = tf.data.Dataset.from_tensor_slices((img_name_train, caption_train))
    dataset = dataset.shuffle(BUFFER_SIZE).batch(BATCH_SIZE).prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
    return dataset
```

```
In [22]: import tensorflow as tf
BATCH_SIZE = 64
BUFFER_SIZE = 1000
train_dataset = create_dataset(x_train,y_train)
test_dataset = create_dataset(x_test,y_test)
```

Modelling

```
In [23]: BATCH_SIZE = 64
BUFFER_SIZE = 1000
embedding_dim = 256
units = 512
vocab_size = len(vocabulary) + 1
num_steps = len(x_train) // BATCH_SIZE
features_shape = 512
attention_features_shape = 49
```

****CNN_ENCODER****

```
In [24]: class VGG16_Encoder(tf.keras.Model):
def __init__(self, embedding_dim):
    super(VGG16_Encoder, self).__init__()
    self.fc = tf.keras.layers.Dense(embedding_dim)
    self.dropout = tf.keras.layers.Dropout(0.5)
def call(self, x):
    x= self.dropout(x)
    x = self.fc(x)
    x = tf.nn.relu(x)
    return x
```

****BahdanauAttention****

```
In [25]: class BahdanauAttention(tf.keras.Model):
def __init__(self, units):
    super(BahdanauAttention, self).__init__()
    self.W1 = tf.keras.layers.Dense(units)
    self.W2 = tf.keras.layers.Dense(units)
    self.V = tf.keras.layers.Dense(1)

def call(self, features, hidden):
    # features(CNN_encoder output) shape == (batch_size, 64, embedding_dim)

    # hidden shape == (batch_size, hidden_size)
    # hidden_with_time_axis shape == (batch_size, 1, hidden_size)
    hidden_with_time_axis = tf.expand_dims(hidden, 1)

    # attention_hidden_layer shape == (batch_size, 64, units)
    attention_hidden_layer = (tf.nn.tanh(self.W1(features) + self.W2(hidden_with_time_axis)))

    # score shape == (batch_size, 64, 1)
    # This gives you an unnormalized score for each image feature.
    score = self.V(attention_hidden_layer)

    # attention_weights shape == (batch_size, 64, 1)
    attention_weights = tf.nn.softmax(score, axis=1)

    # context_vector shape after sum == (batch_size, hidden_size)
    context_vector = attention_weights * features
    context_vector = tf.reduce_sum(context_vector, axis=1)

    return context_vector, attention_weights
```

****Decoder****

```
In [26]: class RNN_Decoder(tf.keras.Model):
def __init__(self, embedding_dim, units, vocab_size):
    super(RNN_Decoder, self).__init__()
    self.units = units
    self.vocab_size=vocab_size
    self.embedding_dim=embedding_dim

    self.embedding = tf.keras.layers.Embedding(self.vocab_size, self.embedding_dim)
    self.gru = tf.keras.layers.GRU(self.units,
                                   return_sequences=True,
                                   return_state=True,
                                   recurrent_initializer='glorot_uniform')

    self.fc1 = tf.keras.layers.Dense(self.units)
    self.fc2 = tf.keras.layers.Dense(self.vocab_size)

    self.attention = BahdanauAttention(self.units)

def call(self, x, features, hidden):
    # defining attention as a separate model
    context_vector, attention_weights = self.attention(features, hidden)

    # x shape after passing through embedding == (64, 1, 256)

    x = self.embedding(x)
    # x shape after concatenation == (64, 1, 256 + 256)
    x = tf.concat([tf.expand_dims(context_vector, 1), x], axis=-1)

    # passing the concatenated vector to the GRU
    output, state = self.gru(x)
    x = self.fc1(output)
    x = tf.reshape(x, (-1, x.shape[2]))
    x = self.fc2(x)

    return x, state, attention_weights

def reset_state(self, batch_size):
    return tf.zeros((batch_size, self.units))
```

```
In [27]: encoder = VGG16_Encoder(embedding_dim)
#vocab_size=4462
decoder = RNN_Decoder(embedding_dim, units, vocab_size)
```

```
In [28]: #custom loss function
optimizer = tf.keras.optimizers.Adam()
loss_object = tf.keras.losses.SparseCategoricalCrossentropy(
    from_logits=True, reduction='none')

def loss_function(real, pred):
    mask = tf.math.logical_not(tf.math.equal(real, 0))
    loss_ = loss_object(real, pred)

    mask = tf.cast(mask, dtype=loss_.dtype)
    loss_ *= mask

    return tf.reduce_mean(loss_)
```

```
In [29]: checkpoint_path = "/content/drive/My Drive/CASE Studies/Image captioning/Attention/train"
ckpt = tf.train.Checkpoint(encoder=encoder,
                           decoder=decoder,
                           optimizer=optimizer)
ckpt_manager = tf.train.CheckpointManager(ckpt, checkpoint_path, max_to_keep=5)

start_epoch = 0
if ckpt_manager.latest_checkpoint:
    start_epoch = int(ckpt_manager.latest_checkpoint.split('-')[-1])
    # restoring the latest checkpoint in checkpoint_path
    ckpt.restore(ckpt_manager.latest_checkpoint)
```

```
In [30]: # Training
loss_plot = []

@tf.function
def train_step(img_tensor, target):
    loss = 0
    hidden = decoder.reset_state(batch_size=target.shape[0])

    dec_input = tf.expand_dims([vocabulary['start']] * target.shape[0], 1)

    with tf.GradientTape() as tape:
        features = encoder(img_tensor)

        for i in range(1, target.shape[1]):
            # passing the features through the decoder
            predictions, hidden, _ = decoder(dec_input, features, hidden)

            loss += loss_function(target[:, i], predictions)

            # using teacher forcing
            dec_input = tf.expand_dims(target[:, i], 1)

    total_loss = (loss / int(target.shape[1]))

    trainable_variables = encoder.trainable_variables + decoder.trainable_variables

    gradients = tape.gradient(loss, trainable_variables)
    optimizer.apply_gradients(zip(gradients, trainable_variables))

    return loss, total_loss
```

```
In [31]: import time
EPOCHS = 40
start_epoch=0
for epoch in range(start_epoch, EPOCHS):
    start = time.time()
    total_loss = 0

    for (batch, (img_tensor, target)) in enumerate(train_dataset):
        batch_loss, t_loss = train_step(img_tensor, target)
        total_loss += t_loss

        if batch % 100 == 0:
            average_batch_loss = batch_loss.numpy()/int(target.shape[1])
            print(f'Epoch {epoch+1} Batch {batch} Loss {average_batch_loss:.4f}
}')
    # storing the epoch end loss value to plot later
    loss_plot.append(total_loss / num_steps)
    if epoch % 5 == 0:
        ckpt_manager.save()
    print(f'Epoch {epoch+1} Loss {total_loss/num_steps:.6f}')
    print(f'Time taken for 1 epoch {time.time()-start:.2f} sec\n')
```

Epoch 1 Batch 0 Loss 2.6503
Epoch 1 Batch 100 Loss 1.5088
Epoch 1 Loss 1.719162
Time taken for 1 epoch 52.17 sec

Epoch 2 Batch 0 Loss 1.5823
Epoch 2 Batch 100 Loss 1.2052
Epoch 2 Loss 1.395903
Time taken for 1 epoch 15.32 sec

Epoch 3 Batch 0 Loss 1.3097
Epoch 3 Batch 100 Loss 1.0028
Epoch 3 Loss 1.199016
Time taken for 1 epoch 14.80 sec

Epoch 4 Batch 0 Loss 1.0878
Epoch 4 Batch 100 Loss 1.0665
Epoch 4 Loss 1.066809
Time taken for 1 epoch 15.24 sec

Epoch 5 Batch 0 Loss 0.9904
Epoch 5 Batch 100 Loss 0.9153
Epoch 5 Loss 0.963886
Time taken for 1 epoch 15.16 sec

Epoch 6 Batch 0 Loss 0.9911
Epoch 6 Batch 100 Loss 0.7892
Epoch 6 Loss 0.881250
Time taken for 1 epoch 15.61 sec

Epoch 7 Batch 0 Loss 0.8170
Epoch 7 Batch 100 Loss 0.7873
Epoch 7 Loss 0.810918
Time taken for 1 epoch 15.35 sec

Epoch 8 Batch 0 Loss 0.7405
Epoch 8 Batch 100 Loss 0.7721
Epoch 8 Loss 0.753207
Time taken for 1 epoch 16.06 sec

Epoch 9 Batch 0 Loss 0.7681
Epoch 9 Batch 100 Loss 0.7041
Epoch 9 Loss 0.703781
Time taken for 1 epoch 15.88 sec

Epoch 10 Batch 0 Loss 0.7247
Epoch 10 Batch 100 Loss 0.5876
Epoch 10 Loss 0.659476
Time taken for 1 epoch 15.40 sec

Epoch 11 Batch 0 Loss 0.5910
Epoch 11 Batch 100 Loss 0.6382
Epoch 11 Loss 0.619931
Time taken for 1 epoch 15.58 sec

Epoch 12 Batch 0 Loss 0.6836
Epoch 12 Batch 100 Loss 0.5133

Epoch 12 Loss 0.585789
Time taken for 1 epoch 15.38 sec

Epoch 13 Batch 0 Loss 0.5868
Epoch 13 Batch 100 Loss 0.4773
Epoch 13 Loss 0.554913
Time taken for 1 epoch 15.14 sec

Epoch 14 Batch 0 Loss 0.5253
Epoch 14 Batch 100 Loss 0.5028
Epoch 14 Loss 0.523101
Time taken for 1 epoch 15.08 sec

Epoch 15 Batch 0 Loss 0.5910
Epoch 15 Batch 100 Loss 0.4918
Epoch 15 Loss 0.499763
Time taken for 1 epoch 15.19 sec

Epoch 16 Batch 0 Loss 0.4266
Epoch 16 Batch 100 Loss 0.5046
Epoch 16 Loss 0.472944
Time taken for 1 epoch 15.56 sec

Epoch 17 Batch 0 Loss 0.4615
Epoch 17 Batch 100 Loss 0.4043
Epoch 17 Loss 0.452070
Time taken for 1 epoch 15.28 sec

Epoch 18 Batch 0 Loss 0.4170
Epoch 18 Batch 100 Loss 0.4546
Epoch 18 Loss 0.429387
Time taken for 1 epoch 15.17 sec

Epoch 19 Batch 0 Loss 0.4680
Epoch 19 Batch 100 Loss 0.4194
Epoch 19 Loss 0.406293
Time taken for 1 epoch 15.20 sec

Epoch 20 Batch 0 Loss 0.3931
Epoch 20 Batch 100 Loss 0.3819
Epoch 20 Loss 0.390194
Time taken for 1 epoch 15.13 sec

Epoch 21 Batch 0 Loss 0.4456
Epoch 21 Batch 100 Loss 0.3311
Epoch 21 Loss 0.373407
Time taken for 1 epoch 15.48 sec

Epoch 22 Batch 0 Loss 0.4193
Epoch 22 Batch 100 Loss 0.3330
Epoch 22 Loss 0.358847
Time taken for 1 epoch 15.24 sec

Epoch 23 Batch 0 Loss 0.3829
Epoch 23 Batch 100 Loss 0.3026
Epoch 23 Loss 0.343873
Time taken for 1 epoch 15.23 sec

Epoch 24 Batch 0 Loss 0.3739
Epoch 24 Batch 100 Loss 0.2983
Epoch 24 Loss 0.329319
Time taken for 1 epoch 15.29 sec

Epoch 25 Batch 0 Loss 0.3714
Epoch 25 Batch 100 Loss 0.2705
Epoch 25 Loss 0.316592
Time taken for 1 epoch 15.12 sec

Epoch 26 Batch 0 Loss 0.3879
Epoch 26 Batch 100 Loss 0.2930
Epoch 26 Loss 0.305256
Time taken for 1 epoch 15.58 sec

Epoch 27 Batch 0 Loss 0.3268
Epoch 27 Batch 100 Loss 0.3047
Epoch 27 Loss 0.292208
Time taken for 1 epoch 15.31 sec

Epoch 28 Batch 0 Loss 0.3481
Epoch 28 Batch 100 Loss 0.2762
Epoch 28 Loss 0.282092
Time taken for 1 epoch 15.32 sec

Epoch 29 Batch 0 Loss 0.3637
Epoch 29 Batch 100 Loss 0.2506
Epoch 29 Loss 0.272510
Time taken for 1 epoch 15.30 sec

Epoch 30 Batch 0 Loss 0.3453
Epoch 30 Batch 100 Loss 0.2494
Epoch 30 Loss 0.263543
Time taken for 1 epoch 15.22 sec

Epoch 31 Batch 0 Loss 0.3078
Epoch 31 Batch 100 Loss 0.2433
Epoch 31 Loss 0.257108
Time taken for 1 epoch 15.49 sec

Epoch 32 Batch 0 Loss 0.3115
Epoch 32 Batch 100 Loss 0.2413
Epoch 32 Loss 0.249230
Time taken for 1 epoch 15.16 sec

Epoch 33 Batch 0 Loss 0.3121
Epoch 33 Batch 100 Loss 0.2498
Epoch 33 Loss 0.241339
Time taken for 1 epoch 15.12 sec

Epoch 34 Batch 0 Loss 0.3202
Epoch 34 Batch 100 Loss 0.1985
Epoch 34 Loss 0.233051
Time taken for 1 epoch 15.12 sec

Epoch 35 Batch 0 Loss 0.2781

```
Epoch 35 Batch 100 Loss 0.2365  
Epoch 35 Loss 0.231847  
Time taken for 1 epoch 15.23 sec
```

```
Epoch 36 Batch 0 Loss 0.2663  
Epoch 36 Batch 100 Loss 0.1989  
Epoch 36 Loss 0.226987  
Time taken for 1 epoch 15.69 sec
```


```
Epoch 37 Batch 0 Loss 0.2968  
Epoch 37 Batch 100 Loss 0.2329  
Epoch 37 Loss 0.220267  
Time taken for 1 epoch 15.28 sec
```

```
Epoch 38 Batch 0 Loss 0.2508  
Epoch 38 Batch 100 Loss 0.2147  
Epoch 38 Loss 0.213822  
Time taken for 1 epoch 15.26 sec
```

```
Epoch 39 Batch 0 Loss 0.2514  
Epoch 39 Batch 100 Loss 0.2263  
Epoch 39 Loss 0.210921  
Time taken for 1 epoch 15.13 sec
```

```
Epoch 40 Batch 0 Loss 0.2926  
Epoch 40 Batch 100 Loss 0.2202  
Epoch 40 Loss 0.208643  
Time taken for 1 epoch 15.12 sec
```

```
In [32]: plt.plot(loss_plot)  
plt.xlabel('Epochs')  
plt.ylabel('Loss')  
plt.title('Loss Plot')  
plt.show()
```



```
In [33]: # creating a dictionary whic is inverse of vocabulary  
invert_vocab={v:k for k ,v in vocabulary.items() }
```

```
In [34]: def evaluate(image):
    attention_plot = np.zeros((max_length, attention_features_shape))

    hidden = decoder.reset_state(batch_size=1)

    img = cv2.imread(image)
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img = cv2.resize(img, (224,224))
    img = np.reshape(img, (1,224,224,3))

    img_tensor_val = image_model.predict(img)#.reshape(1, -1, 512)
    #print(img_tensor_val.shape)
    img_tensor_val = tf.reshape(img_tensor_val, (img_tensor_val.shape[0], -1, img_tensor_val.shape[3]))

    features = encoder(img_tensor_val)

    dec_input = tf.expand_dims([vocabulary['start']], 0)
    result = []

    for i in range(max_length):
        predictions, hidden, attention_weights = decoder(dec_input,
                                                         features,
                                                         hidden)

        attention_plot[i] = tf.reshape(attention_weights, (-1, )).numpy()

        predicted_id = tf.random.categorical(predictions, 1)[0][0].numpy()
        result.append(invert_vocab[predicted_id])

        if invert_vocab[predicted_id] == 'end':
            return result, attention_plot

        dec_input = tf.expand_dims([predicted_id], 0)

    attention_plot = attention_plot[:len(result), :]
    return result, attention_plot
```

```
In [35]: from PIL import Image
def plot_attention(image, result, attention_plot):
    temp_image = np.array(Image.open(image))

    fig = plt.figure(figsize=(10, 10))

    len_result = len(result)
    for i in range(len_result):
        temp_att = np.resize(attention_plot[i], (8, 8))
        grid_size = max(np.ceil(len_result/2), 2)
        ax = fig.add_subplot(grid_size, grid_size, i+1)
        ax.set_title(result[i])
        img = ax.imshow(temp_image)
        ax.imshow(temp_att, cmap='gray', alpha=0.6, extent=img.get_extent())

    plt.tight_layout()
    plt.show()
```

```
In [42]: import nltk.translate.bleu_score as bleu
import warnings
warnings.filterwarnings('ignore')

img=test_images[9]

r,a=evaluate(img)
print("\nGenerated Captions are: ",(" ").join([x for x in r]))
print("\nOriginal captions are :")
key=img.split("/")[1]
score=[]
for i in cap[key]:
    print("\t\t",i)
    reference=i
    translation = (" ").join([x for x in r])
    BLEU_score= bleu.sentence_bleu(reference, translation)
    score.append(BLEU_score)

print("\nThe Best Blue Score is ",max(score))
print(plot_attention(img, r, a))
```

Generated Captions are: a basketball end

Original captions are :

start a man in a field backed by american flags end

start a man is being thrown a frisbee in front of a monument

and flags end

start a man is clapping at the base of the washington memori

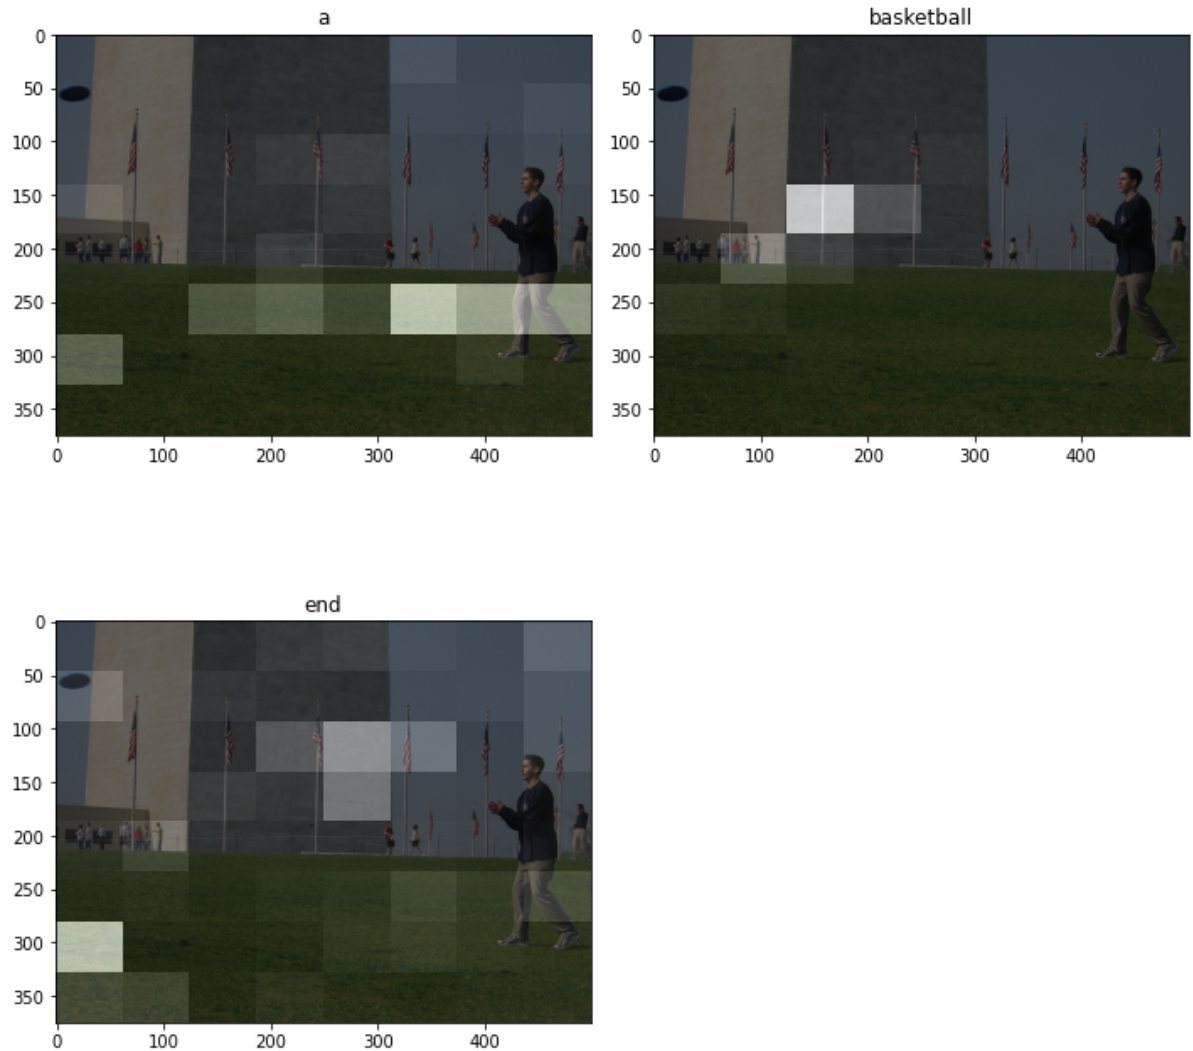
al end

start a man stands in front of the gateway arch end

start a man tries to catch a football on grass surrounded by

american flags end

The Best Blue Score is 0.8891397050194614



None

```
In [37]: import nltk.translate.bleu_score as bleu
import warnings
warnings.filterwarnings('ignore')

img=test_images[0]

r,a=evaluate(img)
print("\nGenerated Captions are: ",(" ").join([x for x in r]))
print("\nOriginal captions are :")
key=img.split("/")[1]
score=[]
for i in cap[key]:
    print("\t\t",i)
    reference=i
    translation = (" ").join([x for x in r])
    BLEU_score= bleu.sentence_bleu(reference, translation)
    score.append(BLEU_score)

print("\nThe Best Blue Score is ",max(score))
print(plot_attention(img, r, a))
```

Generated Captions are: this person doing a helmet end

Original captions are :

start a bicyclist is attempting a trick on a wire in a tropical location end

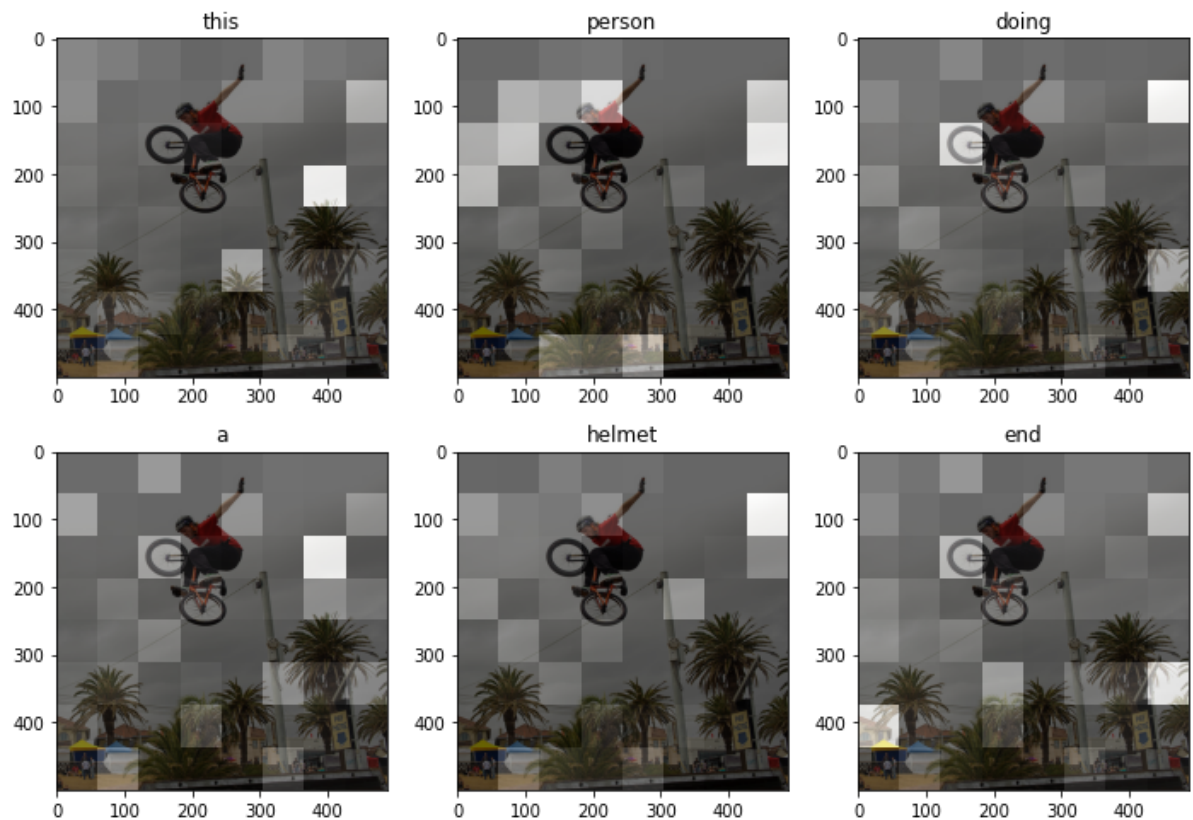
start a biker lets go of the handlebars while flying over the ground end

start an extreme cyclist in red shirt and helmet jumping his bike over a guidewire end

start a person on a bike leaps high in the air at a beach end

start person in orange shirt balancing bicycle on wire end

The Best Blue Score is 0.8408964152537145



None


```
In [38]: import nltk.translate.bleu_score as bleu
import warnings
warnings.filterwarnings('ignore')

img=test_images[40]

r,a=evaluate(img)
print("\nGenerated Captions are: ",(" ").join([x for x in r]))
print("\nOriginal captions are :")
key=img.split("/")[1]
score=[]
for i in cap[key]:
    print("\t\t",i)
    reference=i
    translation = (" ").join([x for x in r])
    BLEU_score= bleu.sentence_bleu(reference, translation)
    score.append(BLEU_score)

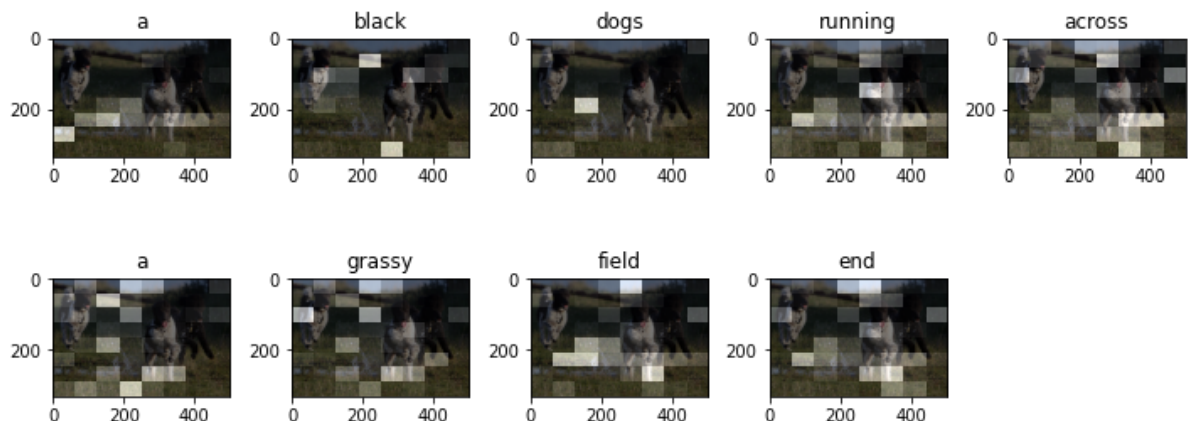
print("\nThe Best Blue Score is ",max(score))
print(plot_attention(img, r, a))
```

Generated Captions are: a black dogs running across a grassy field end

Original captions are :

start the dogs are running through the water end
start three black and white dogs running through a watery fi
eld end
start three dogs running through a puddle end
start three dogs run through a soggy field end
start three dogs run through water and grass end

The Best Blue Score is 0.7796914510717229



None

```
In [46]: import pickle
with open("/content/drive/My Drive/CASE Studies/Image captioning/Attention/vocabul
ary.pkl","wb") as f:
    pickle.dump(vocabulary,f)
```

```
In [47]: import pickle
with open("/content/drive/My Drive/CASE Studies/Image captioning/Attention/invert_vocab.pkl","wb") as f:
    pickle.dump(invert_vocab,f)
```

```
In [ ]:
```