

```
In [1]: from google.colab import drive
drive.mount("/content/drive/")
```

Mounted at /content/drive/

```
In [4]: !wget --header="Host: storage.googleapis.com" --header="User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/90.0.4430.212 Safari/537.36" --header="Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,image/apng,*/*;q=0.8,application/signed-exchange;v=b3;q=0.9" --header="Accept-Language: en-US,en;q=0.9" --header="Referer: https://www.kaggle.com/" "https://storage.googleapis.com/kaggle-data-sets/623289/1111676/bundle/archive.zip?X-Goog-Algorithm=G00G4-RSA-SHA256&X-Goog-Credential=gcp-kaggle-com%40kaggle-161607.iam.gserviceaccount.com%2F20210523%2Fauto%2Fstorage%2Fgoog4_request&X-Goog-Date=20210523T071250Z&X-Goog-Expires=259199&X-Goog-SignedHeaders=host&X-Goog-Signature=4f37286c429b274dd0fca35758b9f0bad4225fa4e8fcff1e366fdcea89b2e79b7587caf3859a1edbc71f7ff10d5a2df38d4eba9e8d77e8c34404d2576f0e2855e65ea1028e9b137bc331ec02b8626160930fdc4c3f0c393ae167ad58e7dc583d3b4d9e96d1f3de0e6700efbb27b16046459567790067b7361df6d4ef1e7daf7152290f09279ac23496514120309e10854ef13ac67978045c454adb58d34d20e5a17352d22f2281d2e5a7e11f3ec433d340a5a85049e7839d0b078c64c77af96983b9fb2fde5a4577770895bee02704f67abc05757f7f21371b06efe19e449e6924f9140bad0d25863b4be820ad37ac1271b0c50bcd6c971c0480bf562ba064d9" -c -O 'archive.zip'
```

```
--2021-05-23 07:13:15-- https://storage.googleapis.com/kaggle-data-sets/623289/1111676/bundle/archive.zip?X-Goog-Algorithm=G00G4-RSA-SHA256&X-Goog-Credential=gcp-kaggle-com%40kaggle-161607.iam.gserviceaccount.com%2F20210523%2Fauto%2Fstorage%2Fgoog4_request&X-Goog-Date=20210523T071250Z&X-Goog-Expires=259199&X-Goog-SignedHeaders=host&X-Goog-Signature=4f37286c429b274dd0fca35758b9f0bad4225fa4e8fcff1e366fdcea89b2e79b7587caf3859a1edbc71f7ff10d5a2df38d4eba9e8d77e8c34404d2576f0e2855e65ea1028e9b137bc331ec02b8626160930fdc4c3f0c393ae167ad58e7dc583d3b4d9e96d1f3de0e6700efbb27b16046459567790067b7361df6d4ef1e7daf7152290f09279ac23496514120309e10854ef13ac67978045c454adb58d34d20e5a17352d22f2281d2e5a7e11f3ec433d340a5a85049e7839d0b078c64c77af96983b9fb2fde5a4577770895bee02704f67abc05757f7f21371b06efe19e449e6924f9140bad0d25863b4be820ad37ac1271b0c50bcd6c971c0480bf562ba064d9
```

Resolving storage.googleapis.com (storage.googleapis.com)... 108.177.119.128, 108.177.126.128, 108.177.127.128, ...

Connecting to storage.googleapis.com (storage.googleapis.com)|108.177.119.128|:443... connected.

HTTP request sent, awaiting response... 200 OK

Length: 1112971163 (1.0G) [application/zip]

Saving to: 'archive.zip'

```
archive.zip          100%[=====>]      1.04G  65.6MB/s   in 18s
```

```
2021-05-23 07:13:34 (57.8 MB/s) - 'archive.zip' saved [1112971163/1112971163]
```



```
In [5]: #Extracting the
from zipfile import ZipFile
with ZipFile("archive.zip","r") as zip:
    zip.extractall()
    print("done")
```


done

Processing The Image

```
In [6]: import matplotlib.pyplot as plt
import cv2
import os
k=os.listdir("Images")
image_list=k[:2000]
image_name=k[:2000]
for i in range(len(image_list)):
    image_list[i]="Images/"+image_list[i]

img=cv2.imread("Images/534200447_b0f3ff02be.jpg")
img=cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
plt.imshow(img)
```

Out[6]: <matplotlib.image.AxesImage at 0x7f20662ffa10>



In [6]:

```
In [8]: # Importing VGG16 for feature extraction from the images
from keras.applications import VGG16 ,ResNet50 ,InceptionV3
vgg=ResNet50(weights="imagenet")
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels.h5
 102973440/102967424 [=====] - 1s 0us/step

```
In [ ]: from keras.models import Model
op = vgg.layers[-2].output
image_model = Model(inputs = vgg.input, outputs = op)
image_model.summary()
```

```
In [11]: # we have to convert all the images to vector of (2048,1)
#for that I m writing a function which takes image and converts it to vector

def img_2_vec(x):
    img=cv2.imread(x)
    img=cv2.resize(img, (224,224))
    img=img.reshape(1,224,224,3)
    img=image_model.predict(img)
    img=img.reshape(2048,)
    return img

img_vector=[]
for i in image_list:
    img_vector.append(img_2_vec(i))
```


```
In [12]: #saving image_name and its vector in dictionary
dict_vector={}
for i in range(len(image_list)):
    dict_vector[image_list[i].split("/")[1]]=img_vector[i]
```

```
In [14]: #selecting only 2000 images
dict_vector=dict([(key ,dict_vector[key]) for key in image_name])
len(dict_vector)
```

```
Out[14]: 2000
```

```
In [15]: # image and its feature vector
img=cv2.imread("Images/534200447_b0f3ff02be.jpg")
img=cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
plt.imshow(img)
print("vector=",dict_vector[image_list[5].split("/")[1]])
```

```
vector= [0.582067    0.71971595 0.04135217 ... 0.776526    0.18743591 0.3457577
]
```



Processing Captions

```
In [16]: # reading the Captions file
with open("captions.txt",'r') as file:
    text=file.readlines()
```

```
In [17]: dict_captions={}
for i in text:
    name=i.split(",")[0]
    caption=i.split(",")[1]
    if name not in dict_captions:
        dict_captions[name]=[caption]
    else:
        dict_captions[name].append(caption)
dict_captions["534200447_b0f3ff02be.jpg"]

cap={}
for i in text:
    name=i.split(",")[0]
    caption=i.split(",")[1]
    if name not in cap:
        cap[name]=[caption]
    else:
        cap[name].append(caption)
```

```
In [18]: dict_captions=dict([(key ,dict_captions[key]) for key in image_name])
len(dict_captions)
cap=dict([(key ,cap[key]) for key in image_name])
```


```
In [19]: #Preprocessing the Captions
import re

def preprocess(d):
    for k,v in d.items():
        for i in range(len(v)):
            v[i]=re.sub('[^A-Za-z]+',' ',v[i])
            v[i]=v[i].lower()
            v[i]=v[i].strip()
            v[i]="start "+v[i]+" end"
    return d
dict_captions=preprocess(dict_captions)
cap=preprocess(cap)
dict_captions['489773343_a8aecf7db3.jpg']
```

```
Out[19]: ['start a small girl in pink and blue end',
'start a young girl in a pink and blue wetsuit and oversized red flip flops
presents a little red ball to a big black dog while another dog looks on en
d',
'start a young girl putting a red ball in a black dogs mouth on the beach en
d',
'start the little girl in the sand is handing a ball to one of the two black
dogs end',
'start two black dogs wait for a little girl to throw a ball at the beach en
d']
```

```
In [20]: #Plotting the image and its captions
import cv2
import matplotlib.pyplot as plt
img=cv2.imread('Images/489773343_a8aecf7db3.jpg')
img=cv2.cvtColor(img,cv2.COLOR_BGR2RGB)
plt.imshow(img)
dict_captions['489773343_a8aecf7db3.jpg']
```

```
Out[20]: ['start a small girl in pink and blue end',
'start a young girl in a pink and blue wetsuit and oversized red flip flops
presents a little red ball to a big black dog while another dog looks on en
d',
'start a young girl putting a red ball in a black dogs mouth on the beach en
d',
'start the little girl in the sand is handing a ball to one of the two black
dogs end',
'start two black dogs wait for a little girl to throw a ball at the beach en
d']
```



```
In [21]: #Creating Vocabulary for the words and giving them the index/integer value
vocab=set()
for k,v in dict_captions.items():
    for i in range(len(v)):
        for j in v[i].split():
            vocab.add(j)
vocab=sorted(vocab)
vocabulary={}
for i,j in enumerate(vocab,1):
    vocabulary[j]=i

print("The Unique Number of words are ",len(vocab))

#trasforming these index to captions
for k,v in dict_captions.items():
    for i in v:
        t=[]
        for j in i.split():
            t.append(vocabulary[j])

        dict_captions[k][v.index(i)]=t

for i in dict_captions['489773343_a8aecf7db3.jpg']:
    print(i)
```

```
The Unique Number of words are 4300
[3569, 1, 3402, 1526, 1846, 2704, 80, 374, 1201]
[3569, 1, 4291, 1526, 1846, 1, 2704, 80, 374, 4178, 80, 2519, 2959, 1374, 139
4, 2815, 1, 2120, 2959, 207, 3862, 1, 333, 354, 1078, 4192, 85, 1078, 2146, 2
462, 1201]
[3569, 1, 4291, 1526, 2886, 1, 2959, 207, 1846, 1, 354, 1079, 2330, 2462, 381
5, 272, 1201]
[3569, 3815, 2120, 1526, 1846, 3815, 3113, 1900, 1659, 1, 207, 3862, 2464, 24
46, 3815, 3997, 354, 1079, 1201]
[3569, 3997, 354, 1079, 4100, 1429, 1, 2120, 1526, 3862, 3836, 1, 207, 146, 3
815, 272, 1201]
```

```
In [23]: #finding the maxlenght of captions from train captions
max_len=[]
for k,v in dict_captions.items():
    for i in v:
        max_len.append(len(i))
max_length=max(max_len)
print("The maximum lenght of caption is",max(max_len))
```

```
The maximum lenght of caption is 33
```

Modelling

```
In [24]: #splitting the according to images names
from sklearn.model_selection import train_test_split
train,test=train_test_split(image_name,test_size=0.2)
tr,cv=train_test_split(train,test_size=0.2)
```

```
In [25]: # this function will be use
from keras.utils import to_categorical
from keras.preprocessing.sequence import pad_sequences
import numpy as np

VOCAB_SIZE = len(vocabulary)

def generator(image, caption):
    X = []
    y_in = []
    y_out = []
    for k, v in caption.items():
        for j in v:
            for i in range(1, len(j)):
                X.append(image[k])
                in_seq= [j[:i]]
                out_seq = [j[i]]
                in_seq = pad_sequences(in_seq, maxlen=max_length, padding='post', truncating='post')[0]
                out_seq = to_categorical([out_seq], num_classes=VOCAB_SIZE+1)[0]
                y_in.append(in_seq)
                y_out.append(out_seq)

    X=np.array(X)
    y_in=np.array(y_in)
    y_out=np.array(y_out)

    return X,y_in,y_out
```

```
In [26]: #dict_captions.pop('image')
train_vector=dict([(key, dict_vector[key]) for key in tr])
train_captions=dict([(key, dict_captions[key]) for key in tr])

cv_vector=dict([(key, dict_vector[key]) for key in cv])
cv_captions=dict([(key, dict_captions[key]) for key in cv])

test_vector=dict([(key, dict_vector[key]) for key in test])
test_captions=dict([(key, dict_captions[key]) for key in test])
```

```
In [27]: #dict_captions.pop('image')
train_x,train_in,train_out=generator(train_vector,train_captions)
cv_x,cv_in,cv_out=generator(cv_vector,cv_captions)
test_x,test_in,test_out=generator(test_vector,test_captions)
```



```
In [28]: import pickle
         embeddings_index = dict()
         with open('/content/drive/My Drive/glove_vectors', 'rb') as f:
             model = pickle.load(f)
             glove_words = set(model.keys())

         essay_vocab = len(vocabulary) + 1
         essay_matrix = np.zeros((essay_vocab, 300))
         for word, i in vocabulary.items():
             embedding_vector = model.get(word)
             if embedding_vector is not None:
                 essay_matrix[i] = embedding_vector
```

```
In [29]: essay_matrix.shape
```

```
Out[29]: (4301, 300)
```

```

In [ ]: # model buliding
from keras.models import Model, Sequential
from keras.layers import Input, Dense, Embedding, LSTM, Dropout, add, Concatenate, BatchNormalization, Flatten, RepeatVector, TimeDistributed, Activation

from keras import regularizers

inputs1 = Input(shape=(2048,))
dense = Dense(512, activation='relu', kernel_initializer="HeUniform")(inputs1)
D1 = Dropout(0.5)(dense)
B1=BatchNormalization()(D1)
dense = Dense(256, activation='relu', kernel_initializer="HeUniform")(B1)
image_out=RepeatVector(max_length)(dense)

inputs2 = Input(shape=(max_length,))
E1 = Embedding(len(vocabulary)+1, 300, weights=[essay_matrix], trainable=False)(inputs2)
D2 = Dropout(0.5)(E1)
lstm = LSTM(256, return_sequences=True)(D2)
caption_out=TimeDistributed(Dense(256))(lstm)

concat = Concatenate()([image_out, caption_out])
x = LSTM(128, return_sequences=True)(concat)
x = LSTM(512, return_sequences=False)(x)
dense = Dense(256, activation='relu', kernel_initializer="HeUniform")(x)
D3=Dropout(0.5)(dense)
B3=BatchNormalization()(D3)
dense = Dense(256, activation='relu', kernel_initializer="HeUniform")(B3)
outputs = Dense(len(vocabulary)+1, activation='softmax')(dense)
model = Model(inputs=[inputs1, inputs2], outputs=outputs)
model.summary()

```

Model: "model_8"

Layer (type)	Output Shape	Param #	Connected to
input_20 (InputLayer)	[(None, 2048)]	0	
dense_36 (Dense) [0]	(None, 512)	1049088	input_20[0]
input_21 (InputLayer)	[(None, 33)]	0	
dropout_25 (Dropout) [0]	(None, 512)	0	dense_36[0]
embedding_9 (Embedding) [0]	(None, 33, 300)	1290300	input_21[0]
batch_normalization_16 (Batch Normalization) [0][0]	(None, 512)	2048	dropout_25
dropout_26 (Dropout) [0][0]	(None, 33, 300)	0	embedding_9
dense_37 (Dense) batch_normalization_16[0][0]	(None, 256)	131328	batch_normalization_16[0][0]
lstm_12 (LSTM) [0][0]	(None, 33, 256)	570368	dropout_26
repeat_vector_3 (RepeatVector) [0]	(None, 33, 256)	0	dense_37[0]
time_distributed_3 (TimeDistributed) [0]	(None, 33, 256)	65792	lstm_12[0]
concatenate_8 (Concatenate) repeat_vector_3[0][0] time_distributed_3[0][0]	(None, 33, 512)	0	repeat_vector_3[0] time_distributed_3[0][0]
lstm_13 (LSTM) concatenate_8[0][0]	(None, 33, 128)	328192	concatenate_8[0][0]

lstm_14 (LSTM) [0]	(None, 512)	1312768	lstm_13[0]
dense_39 (Dense) [0]	(None, 256)	131328	lstm_14[0]
dropout_27 (Dropout) [0]	(None, 256)	0	dense_39[0]
batch_normalization_17 (BatchNo [0])[0]	(None, 256)	1024	dropout_27
dense_40 (Dense) ization_17[0][0]	(None, 256)	65792	batch_normal
dense_41 (Dense) [0]	(None, 4301)	1105357	dense_40[0]

=====
=====
Total params: 6,053,385
Trainable params: 4,761,549
Non-trainable params: 1,291,836



```
In [ ]: model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=["accuracy"])
        model.fit([train_x, train_in], train_out, batch_size=512, epochs=20, validation_data=([cv_x, cv_in], cv_out))
```

Epoch 1/20
144/144 [=====] - 19s 102ms/step - loss: 5.9207 - accuracy: 0.1296 - val_loss: 5.1523 - val_accuracy: 0.1933

Epoch 2/20
144/144 [=====] - 14s 96ms/step - loss: 4.7055 - accuracy: 0.1926 - val_loss: 4.8256 - val_accuracy: 0.1945

Epoch 3/20
144/144 [=====] - 14s 97ms/step - loss: 4.5215 - accuracy: 0.2036 - val_loss: 4.6425 - val_accuracy: 0.2322

Epoch 4/20
144/144 [=====] - 14s 98ms/step - loss: 4.3334 - accuracy: 0.2349 - val_loss: 4.4816 - val_accuracy: 0.2603

Epoch 5/20
144/144 [=====] - 14s 100ms/step - loss: 4.0407 - accuracy: 0.2634 - val_loss: 4.2341 - val_accuracy: 0.2817

Epoch 6/20
144/144 [=====] - 14s 100ms/step - loss: 3.8424 - accuracy: 0.2825 - val_loss: 4.1155 - val_accuracy: 0.2968

Epoch 7/20
144/144 [=====] - 14s 101ms/step - loss: 3.7120 - accuracy: 0.2956 - val_loss: 4.0413 - val_accuracy: 0.3034

Epoch 8/20
144/144 [=====] - 15s 101ms/step - loss: 3.6170 - accuracy: 0.3002 - val_loss: 4.1324 - val_accuracy: 0.2975

Epoch 9/20
144/144 [=====] - 15s 102ms/step - loss: 3.5162 - accuracy: 0.3063 - val_loss: 3.9819 - val_accuracy: 0.3057

Epoch 10/20
144/144 [=====] - 15s 103ms/step - loss: 3.4322 - accuracy: 0.3136 - val_loss: 3.9601 - val_accuracy: 0.3135

Epoch 11/20
144/144 [=====] - 15s 103ms/step - loss: 3.3343 - accuracy: 0.3231 - val_loss: 3.9231 - val_accuracy: 0.3233

Epoch 12/20
144/144 [=====] - 15s 102ms/step - loss: 3.2772 - accuracy: 0.3269 - val_loss: 3.8929 - val_accuracy: 0.3290

Epoch 13/20
144/144 [=====] - 15s 102ms/step - loss: 3.2133 - accuracy: 0.3341 - val_loss: 3.8910 - val_accuracy: 0.3275

Epoch 14/20
144/144 [=====] - 15s 102ms/step - loss: 3.1658 - accuracy: 0.3356 - val_loss: 3.9050 - val_accuracy: 0.3288

Epoch 15/20
144/144 [=====] - 15s 102ms/step - loss: 3.1108 - accuracy: 0.3390 - val_loss: 3.8884 - val_accuracy: 0.3298

Epoch 16/20
144/144 [=====] - 15s 102ms/step - loss: 3.0575 - accuracy: 0.3438 - val_loss: 3.9162 - val_accuracy: 0.3274

Epoch 17/20
144/144 [=====] - 15s 103ms/step - loss: 2.9896 - accuracy: 0.3450 - val_loss: 3.8932 - val_accuracy: 0.3326

Epoch 18/20
144/144 [=====] - 15s 103ms/step - loss: 2.9552 - accuracy: 0.3517 - val_loss: 3.9078 - val_accuracy: 0.3361

Epoch 19/20
144/144 [=====] - 15s 103ms/step - loss: 2.9281 - accuracy: 0.3548 - val_loss: 3.8705 - val_accuracy: 0.3373

Epoch 20/20

144/144 [=====] - 15s 103ms/step - loss: 2.8811 - accuracy: 0.3580 - val_loss: 3.9119 - val_accuracy: 0.3370

Out[]: <tensorflow.python.keras.callbacks.History at 0x7fd8414872d0>

```
In [ ]: import tensorflow as tf
model.compile(loss='categorical_crossentropy', optimizer=tf.keras.optimizers.Adam(learning_rate=0.0001), metrics=["accuracy"])
model.fit([train_x, train_in], train_out, batch_size=512, epochs=20, validation_data=([test_x, test_in], test_out))
```


Epoch 1/20
144/144 [=====] - 20s 111ms/step - loss: 2.7509 - accuracy: 0.3741 - val_loss: 3.6807 - val_accuracy: 0.3535

Epoch 2/20
144/144 [=====] - 15s 105ms/step - loss: 2.7209 - accuracy: 0.3760 - val_loss: 3.6966 - val_accuracy: 0.3535

Epoch 3/20
144/144 [=====] - 15s 106ms/step - loss: 2.6930 - accuracy: 0.3814 - val_loss: 3.7154 - val_accuracy: 0.3534

Epoch 4/20
144/144 [=====] - 15s 106ms/step - loss: 2.6710 - accuracy: 0.3810 - val_loss: 3.7276 - val_accuracy: 0.3550

Epoch 5/20
144/144 [=====] - 15s 105ms/step - loss: 2.6635 - accuracy: 0.3816 - val_loss: 3.7452 - val_accuracy: 0.3562

Epoch 6/20
144/144 [=====] - 15s 105ms/step - loss: 2.6334 - accuracy: 0.3850 - val_loss: 3.7660 - val_accuracy: 0.3550

Epoch 7/20
144/144 [=====] - 15s 105ms/step - loss: 2.6357 - accuracy: 0.3856 - val_loss: 3.7781 - val_accuracy: 0.3545

Epoch 8/20
144/144 [=====] - 15s 105ms/step - loss: 2.6239 - accuracy: 0.3896 - val_loss: 3.7909 - val_accuracy: 0.3567

Epoch 9/20
144/144 [=====] - 15s 105ms/step - loss: 2.6122 - accuracy: 0.3859 - val_loss: 3.8144 - val_accuracy: 0.3544

Epoch 10/20
144/144 [=====] - 15s 105ms/step - loss: 2.6023 - accuracy: 0.3890 - val_loss: 3.8230 - val_accuracy: 0.3576

Epoch 11/20
144/144 [=====] - 15s 105ms/step - loss: 2.5888 - accuracy: 0.3898 - val_loss: 3.8370 - val_accuracy: 0.3558

Epoch 12/20
144/144 [=====] - 15s 105ms/step - loss: 2.5789 - accuracy: 0.3946 - val_loss: 3.8482 - val_accuracy: 0.3543

Epoch 13/20
144/144 [=====] - 15s 105ms/step - loss: 2.5729 - accuracy: 0.3913 - val_loss: 3.8602 - val_accuracy: 0.3568

Epoch 14/20
144/144 [=====] - 15s 105ms/step - loss: 2.5631 - accuracy: 0.3934 - val_loss: 3.8824 - val_accuracy: 0.3535

Epoch 15/20
144/144 [=====] - 15s 105ms/step - loss: 2.5485 - accuracy: 0.3930 - val_loss: 3.8839 - val_accuracy: 0.3559

Epoch 16/20
144/144 [=====] - 15s 105ms/step - loss: 2.5436 - accuracy: 0.3975 - val_loss: 3.9005 - val_accuracy: 0.3557

Epoch 17/20
144/144 [=====] - 15s 105ms/step - loss: 2.5487 - accuracy: 0.3959 - val_loss: 3.9064 - val_accuracy: 0.3556

Epoch 18/20
144/144 [=====] - 15s 105ms/step - loss: 2.5279 - accuracy: 0.3966 - val_loss: 3.9186 - val_accuracy: 0.3555

Epoch 19/20
144/144 [=====] - 15s 106ms/step - loss: 2.5143 - accuracy: 0.3996 - val_loss: 3.9353 - val_accuracy: 0.3549

Epoch 20/20

144/144 [=====] - 15s 105ms/step - loss: 2.5146 - accuracy: 0.3997 - val_loss: 3.9462 - val_accuracy: 0.3547

Out[]: <tensorflow.python.keras.callbacks.History at 0x7fd83fb03e50>

```
In [33]: from keras.models import load_model
#model.save('/content/drive/My Drive/CASE Studies/Image captioning/model.h5')
model=load_model("/content/drive/My Drive/CASE Studies/Image captioning/model.h5")
```

Prediction

```
In [34]: # creating a dictionary whic is inverse of vocabulary
invert_vocab={v:k for k ,v in vocabulary.items()}
```

```
In [35]: #funtion for making single prediction
def prediction(img):
    img = cv2.imread(img)
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
    img = cv2.resize(img, (224,224))
    plt.imshow(img)
    img = np.reshape(img, (1,224,224,3))

    x=image_model.predict(img).reshape(1,2048)

    y_in=["start"]
    caption=''

    for i in range(max_length):
        t=[]
        for j in y_in:
            t.append(vocabulary[j])

        input=pad_sequences([t], maxlen=max_length, padding='post', truncating='post')
        pred=np.argmax(model.predict([x,input]))

        caption=caption + " " +invert_vocab[pred]

        if invert_vocab[pred]=="end":
            break

        y_in.append(invert_vocab[pred])
    return caption
```

```
In [46]: import nltk.translate.bleu_score as bleu
import warnings
warnings.filterwarnings('ignore')

img="Images/"+test[71]
caption=prediction(img)
caption=caption.replace("end","")
print("\n The caption generated is==>",caption)
print("\n The original captions are:" )
score=[]
for i in cap[test[71]]:
    print("\t\t",i)
    reference=i
    translation = caption
    BLEU_score= bleu.sentence_bleu(reference, translation)
    score.append(BLEU_score)


print("\nThe Best  Blue Score is ",max(score))
```

The caption generated is==> a boy is climbing a rock

The original captions are:

```
start a climber is climbing on an artificial climbing wall e
nd
start a man attempting to climb a fake rock wall end
start a man climbs a rock wall end
start a man climbs up an artificial climbing wall end
start a rock climber climbs a rock wall end
```

The Best Blue Score is 0.8408964152537145



```
In [45]: import nltk.translate.bleu_score as bleu
import warnings
warnings.filterwarnings('ignore')

img="Images/"+test[48]
caption=prediction(img)
print("\n The caption generated is==>",caption)
print("\n The original captions are:" )
score=[]
for i in cap[test[48]]:
    print("\t\t",i)
    reference=i
    translation = caption
    BLEU_score= bleu.sentence_bleu(reference, translation)
    score.append(BLEU_score)

print("\nThe Best  Blue Score is ",max(score))
```

The caption generated is==> a dog is jumping over a stick end

The original captions are:

start a black dog is leaping into a creek end

start a black dog is leaping over a stream of running water

end

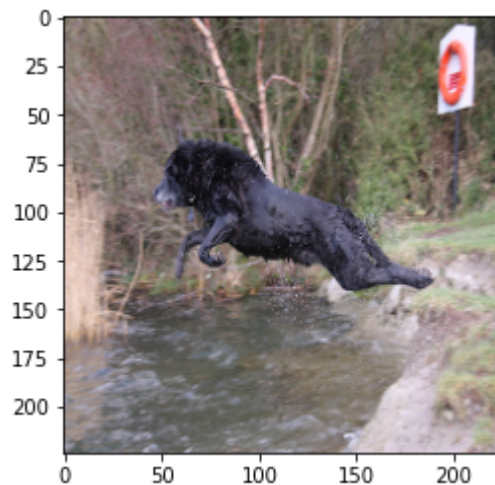
start a black dog jumping off a river bank near a wooded are

a end

start a wet black dog jumping into a lake end

start a wet black dog leaps across a stream end

The Best Blue Score is 0.8529987544592307



```
In [38]: import nltk.translate.bleu_score as bleu
import warnings
warnings.filterwarnings('ignore')

img="Images/"+test[187]
caption=prediction(img)
print("\n The caption generated is==>",caption)
print("\n The original captions are:" )
score=[]
for i in cap[test[187]]:
    print("\t\t",i)
    reference=i
    translation = caption
    BLEU_score= bleu.sentence_bleu(reference, translation)
    score.append(BLEU_score)

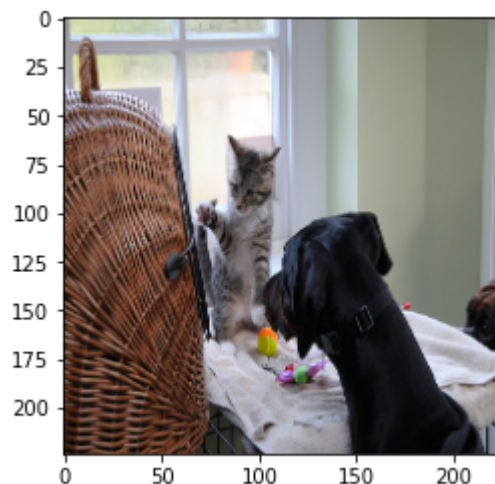
print("\nThe Best  Blue Score is ",max(score))
```

The caption generated is==> a dog is laying on a bed end

The original captions are:

```
start a black dog is looking at the tabby cat that is sittin
g by the window end
start a dog and cat sit by a table with toys on it end
start a tawny cat is playing with a basket while a black dog
watches end
start big black dog and kitten curiously looking at the wood
en weave basket end
start the pets play in the apartment end
```

The Best Blue Score is 0.8020396005825877



```
In [ ]: import pickle
with open("/content/drive/My Drive/CASE Studies/Image captioning/vocabulary.pk
l", "wb") as f:
    pickle.dump(vocabulary,f)
```

```
In [ ]: import pickle
with open("/content/drive/My Drive/CASE Studies/Image captioning/invert_vocabulary.pkl","wb") as f:
    pickle.dump(invert_vocab,f)
```

```
In [ ]: import nltk.translate.bleu_score as bleu
import warnings
warnings.filterwarnings('ignore')
s=[]
for j in test:
    img="Images/"+j
    caption=prediction(img)
    score=[]
    for i in cap[j]:
        reference=i
        translation = caption
        BLEU_score= bleu.sentence_bleu(reference, translation)
        score.append(BLEU_score)

    s.append(max(score))
```

```
In [40]: s=np.array(s)
print("The Median BleuScore is ",np.median(s))
print("The Mean BleuScore is ",s.mean())
print("The Max BleuScore is ",max(s))
print("The Min BleuScore is ",min(s))
```

```
The Median BleuScore is  0.7561559858236034
The Mean BleuScore is  0.7328405376797648
The Max BleuScore is  0.8907817059279295
The Min BleuScore is  0.3932836841548814
```

```
In [ ]:
```